

Mobile Price Classification

Bob has started his own mobile company. He wants to give tough fight to big companies like Apple,Samsung etc.

He does not know how to estimate price of mobiles his company creates. In this competitive mobile phone market you cannot simply assume things. To solve this problem he collects sales data of mobile phones of various companies.

Bob wants to find out some relation between features of a mobile phone(eg:- RAM,Internal Memory etc) and its selling price. But he is not so good at Machine Learning. So he needs your help to solve this problem.

In this problem you do not have to predict actual price but a price range indicating how high the price is

Let's Dive into it

Import necessary libraries

```
In [5]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

Read 'mobile_price.csv' dataset and store it in a DataFrame

```
In [6]: df = pd.read_csv('mobile_price.csv')
```

View top 5 rows

```
In [7]: pd.set_option('display.max_columns', None) # Used to view non-truncated columns
```

```
In [8]: df.head()
```

```
Out[8]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_co
0	842	0	2.2	0	1	0		7	0.6	188
1	1021	1	0.5	1	0	1		53	0.7	136
2	563	1	0.5	1	2	1		41	0.9	145
3	615	1	2.5	0	0	0		10	0.8	131
4	1821	1	1.2	0	13	1		44	0.6	141



View info of the dataset

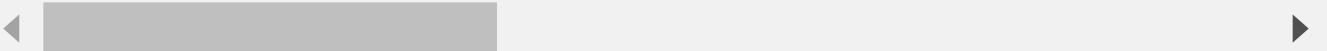
In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   battery_power    2000 non-null   int64  
 1   blue              2000 non-null   int64  
 2   clock_speed      2000 non-null   float64 
 3   dual_sim          2000 non-null   int64  
 4   fc                2000 non-null   int64  
 5   four_g            2000 non-null   int64  
 6   int_memory        2000 non-null   int64  
 7   m_dep             2000 non-null   float64 
 8   mobile_wt         2000 non-null   int64  
 9   n_cores           2000 non-null   int64  
 10  pc                2000 non-null   int64  
 11  px_height         2000 non-null   int64  
 12  px_width          2000 non-null   int64  
 13  ram               2000 non-null   int64  
 14  sc_h              2000 non-null   int64  
 15  sc_w              2000 non-null   int64  
 16  talk_time          2000 non-null   int64  
 17  three_g            2000 non-null   int64  
 18  touch_screen       2000 non-null   int64  
 19  wifi               2000 non-null   int64  
 20  price_range        2000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

View basic statistical information about the dataset

In [10]: `df.describe()`

Out[10]:	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.046500
std	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.145715
min	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000
25%	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000
max	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000



Check for null values

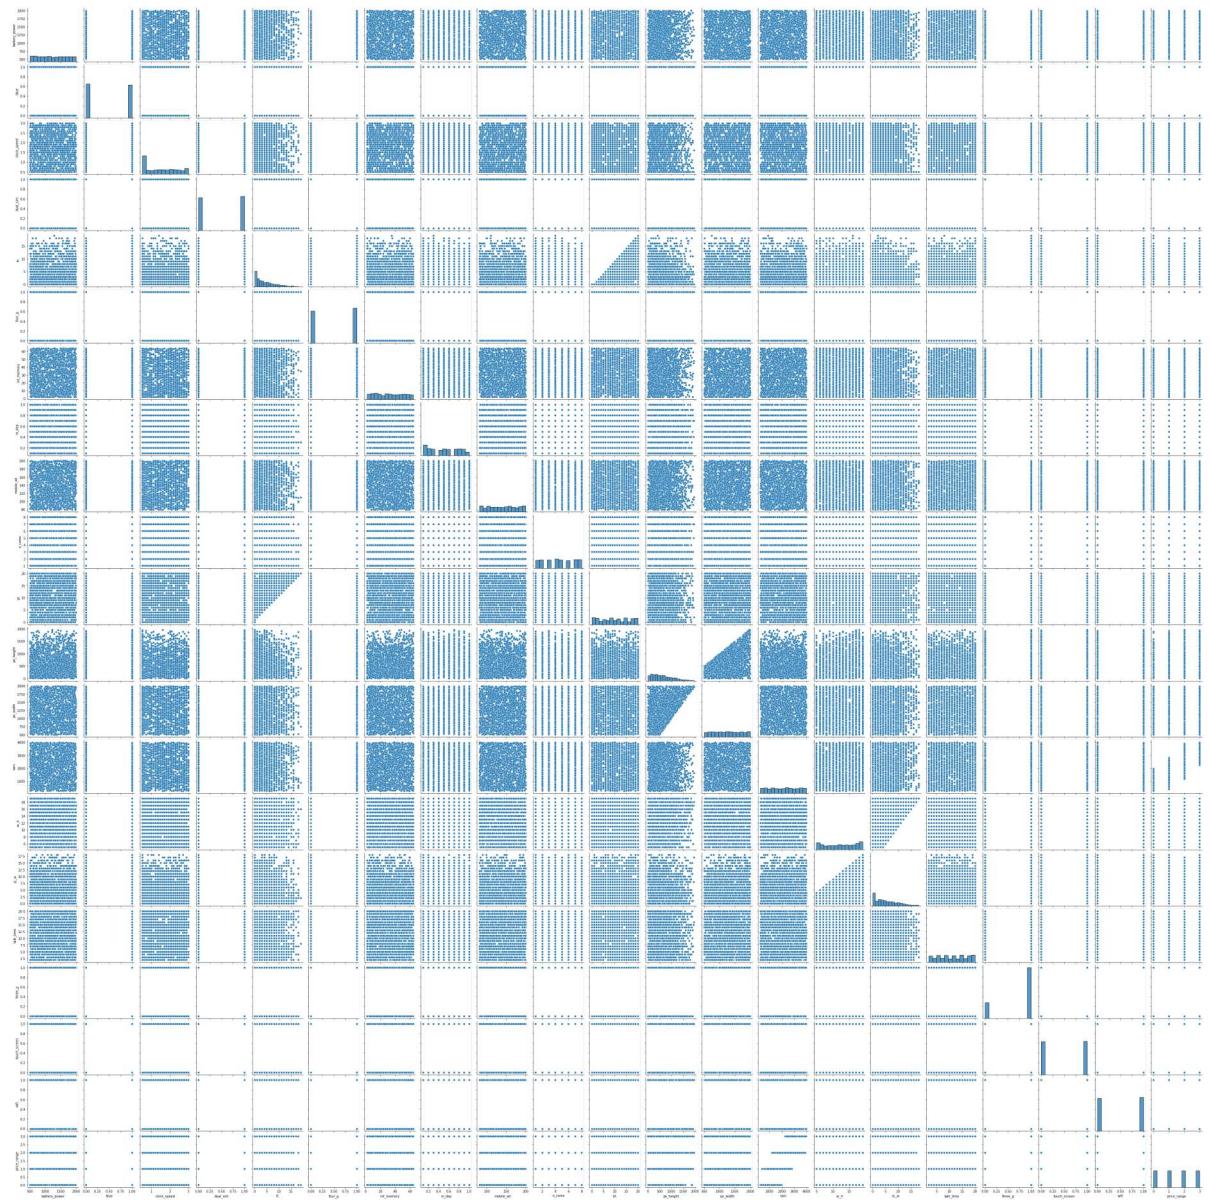
In [7]: `df.isna().sum()`

```
Out[7]: battery_power      0  
        blue                0  
        clock_speed         0  
        dual_sim            0  
        fc                  0  
        four_g              0  
        int_memory          0  
        m_dep               0  
        mobile_wt            0  
        n_cores              0  
        pc                  0  
        px_height            0  
        px_width             0  
        ram                 0  
        sc_h                0  
        sc_w                0  
        talk_time            0  
        three_g              0  
        touch_screen          0  
        wifi                 0  
        price_range           0  
        dtype: int64
```

Visualization

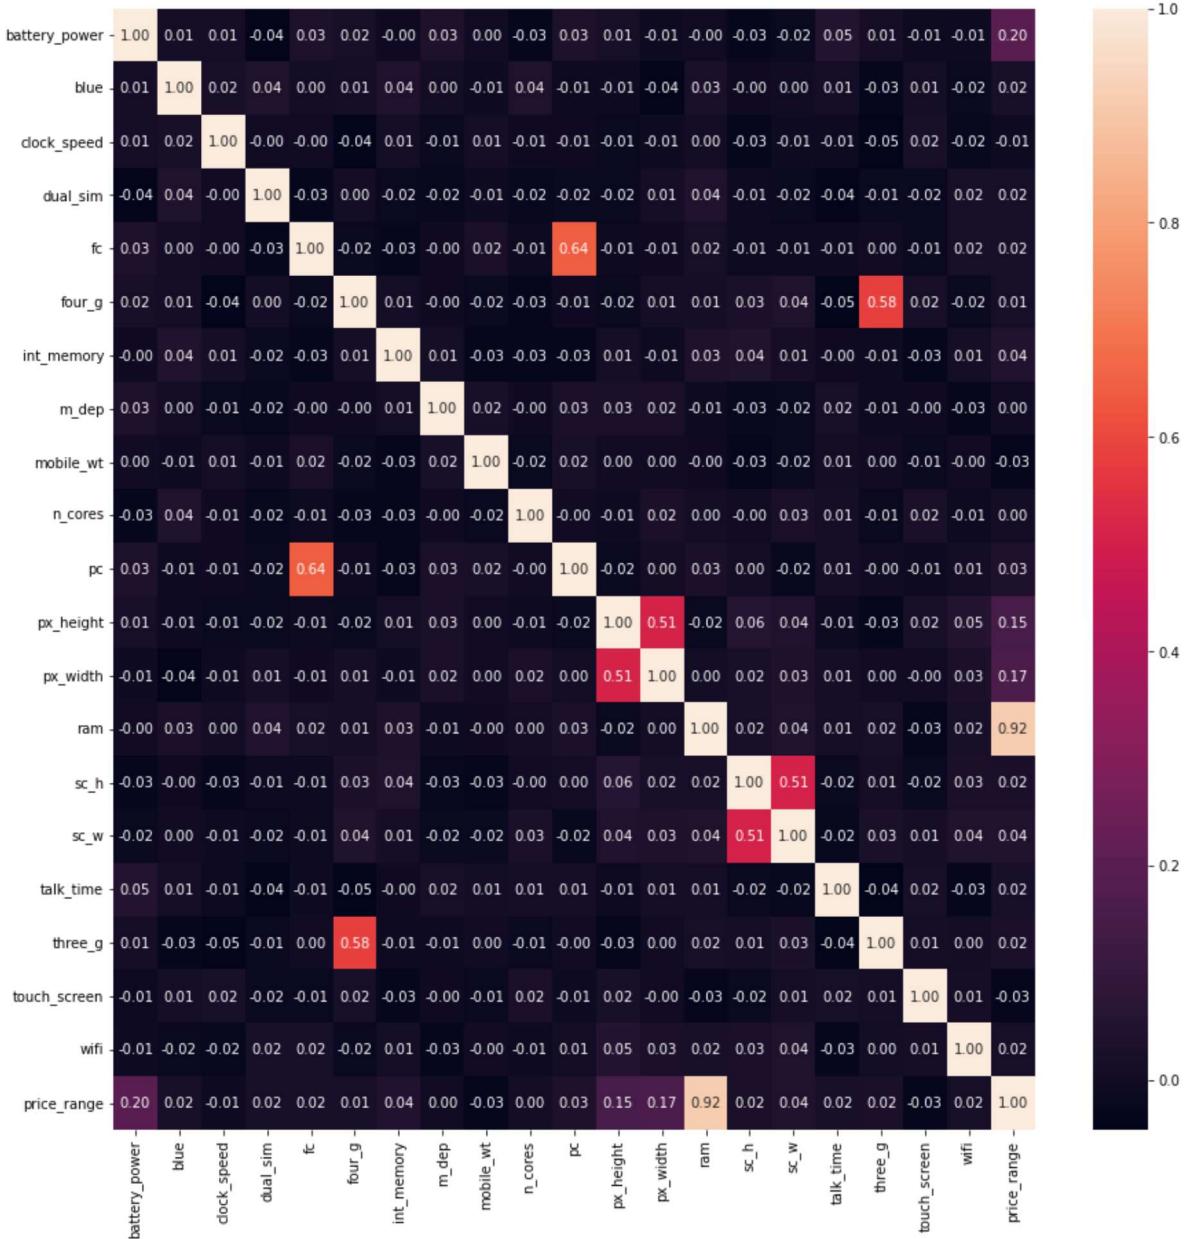
Plot a pairplot of the dataset

```
In [14]: sns.pairplot(df)  
plt.show()
```



Plot a heatmap to show the correlation between features

```
In [17]: plt.figure(figsize = (15,15))
sns.heatmap(df.corr(), annot=True, fmt='.2f')
plt.show()
```



Split the data into input and target variables

```
In [11]: X = df.drop(columns = ['price_range'])
y = df['price_range']
```

Standardise the data with StandardScaler

```
In [12]: from sklearn.preprocessing import StandardScaler
```

```
In [13]: scaler = StandardScaler()
```

```
In [14]: xcolumns = X.columns
```

```
In [15]: X = scaler.fit_transform(X)
```

```
In [13]: X = pd.DataFrame(X, columns = xcolumns)
```

```
In [14]: X.head()
```

Out[14]:	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep
0	-0.902597	-0.990050	0.830779	-1.019184	-0.762495	-1.043966	-1.380644	0.340740
1	-0.495139	1.010051	-1.253064	0.981177	-0.992890	0.957886	1.155024	0.687548
2	-1.537686	1.010051	-1.253064	0.981177	-0.532099	0.957886	0.493546	1.381165
3	-1.419319	1.010051	1.198517	-1.019184	-0.992890	-1.043966	-1.215274	1.034357
4	1.325906	1.010051	-0.395011	-1.019184	2.002254	0.957886	0.658915	0.340740

◀ ▶

Split the dataset into training and testing set

```
In [21]: from sklearn.model_selection import train_test_split
```

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.25, random_
```

```
In [24]:
```

```
In [24]: X_train.shape
```

```
Out[24]: (1500, 20)
```

```
In [25]: X_test.shape
```

```
Out[25]: (500, 20)
```

Create random forest model.

Let n_estimator = 100

```
In [26]: from sklearn.ensemble import RandomForestClassifier
```

```
In [34]: model = RandomForestClassifier(n_estimators= 100)
```

Train the model

```
In [35]: model.fit(X_train, y_train)
```

```
Out[35]: RandomForestClassifier()
```

Check the score of our model

```
In [36]: model.score(X_train, y_train)
```

```
Out[36]: 1.0
```

Predict using X_test

```
In [37]: y_pred = model.predict(X_test)
y_pred
```

```
Out[37]: array([0, 1, 1, 3, 2, 3, 3, 2, 3, 1, 0, 1, 3, 3, 0, 2, 2, 3, 2, 0, 0, 3,
   1, 0, 0, 1, 3, 0, 1, 2, 0, 3, 0, 1, 0, 3, 3, 1, 1, 3, 3, 3, 0, 0,
   3, 3, 2, 0, 1, 2, 3, 2, 3, 2, 1, 3, 1, 3, 1, 0, 3, 2, 3, 0, 2, 0,
   0, 3, 3, 2, 0, 1, 2, 0, 2, 3, 0, 3, 2, 3, 1, 0, 3, 0, 2, 2, 1, 2,
   3, 3, 2, 3, 3, 2, 0, 2, 3, 1, 1, 1, 0, 0, 3, 3, 2, 0, 0, 1, 1, 1,
   0, 1, 3, 2, 0, 3, 1, 3, 0, 1, 3, 1, 0, 1, 3, 2, 3, 2, 0, 3, 1, 0,
   1, 2, 3, 1, 3, 2, 0, 0, 3, 1, 3, 3, 1, 2, 2, 2, 0, 2, 0, 3, 1, 2,
   1, 0, 2, 3, 0, 1, 3, 2, 3, 0, 1, 3, 0, 3, 3, 0, 3, 1, 2, 1, 1, 1,
   2, 1, 3, 1, 0, 2, 1, 0, 0, 3, 1, 2, 3, 2, 2, 1, 2, 1, 2, 2, 3, 1,
   2, 1, 3, 1, 2, 1, 3, 0, 1, 3, 3, 2, 0, 1, 2, 0, 0, 2, 3, 0, 1, 3,
   0, 3, 0, 0, 1, 1, 2, 3, 2, 0, 1, 2, 0, 1, 1, 3, 3, 2, 0, 1, 3,
   2, 1, 0, 3, 2, 1, 2, 1, 1, 2, 3, 1, 3, 2, 1, 1, 2, 3, 1, 3, 2, 2,
   2, 2, 3, 1, 2, 1, 0, 3, 2, 0, 2, 1, 1, 3, 2, 2, 2, 1, 3, 1, 0, 2,
   3, 0, 0, 1, 2, 2, 1, 2, 2, 0, 2, 3, 0, 3, 3, 0, 0, 3, 0, 3, 0,
   0, 0, 3, 0, 1, 2, 0, 1, 1, 3, 1, 2, 3, 3, 0, 0, 0, 0, 0, 2, 1, 1,
   3, 1, 0, 0, 1, 2, 3, 1, 2, 1, 0, 2, 0, 0, 3, 2, 1, 1, 0, 2, 0, 3,
   3, 1, 2, 3, 0, 0, 1, 2, 3, 3, 1, 0, 3, 1, 2, 0, 0, 2, 0, 0, 2,
   2, 3, 0, 1, 1, 3, 0, 1, 0, 1, 0, 1, 0, 2, 2, 0, 1, 3, 0, 1, 0, 3,
   1, 3, 0, 0, 0, 0, 0, 3, 0, 0, 3, 3, 1, 1, 3, 0, 0, 2, 1, 0, 0, 1, 2,
   0, 0, 3, 0, 0, 0, 2, 1, 1, 1, 0, 3, 0, 0, 2, 3, 1, 0, 2, 1, 3, 1,
   0, 0, 0, 0, 2, 0, 1, 1, 0, 3, 3, 2, 2, 3, 0, 3, 2, 3, 3, 0, 0, 2,
   0, 1, 1, 0, 0, 2, 0, 2, 3, 0, 1, 1, 1, 1, 2, 1, 2, 1, 0, 2, 3, 2,
   2, 1, 0, 3, 1, 1, 0, 2, 2, 2, 0, 1, 1, 2, 2, 2, 0, 1, 2, 2, 2], dtype=int64)
```

Check the accuracy score of our prediction

```
In [38]: from sklearn import metrics
```

```
In [39]: metrics.accuracy_score(y_test, y_pred)
```

```
Out[39]: 0.892
```

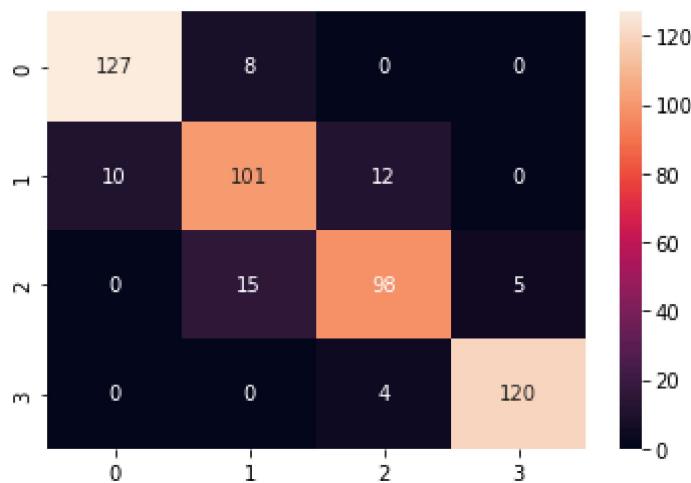
Create a confusion matrix

```
In [40]: metrics.confusion_matrix(y_test, y_pred)
```

```
Out[40]: array([[127,    8,    0,    0],
   [ 10, 101,   12,    0],
   [  0,  15,  98,    5],
   [  0,    0,   4, 120]], dtype=int64)
```

Plot confusion matrix on heatmap

```
In [41]: sns.heatmap(metrics.confusion_matrix(y_test,y_pred), annot = True, fmt = 'd')
plt.show()
```



Create classification report

```
In [42]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.93	0.94	0.93	135
1	0.81	0.82	0.82	123
2	0.86	0.83	0.84	118
3	0.96	0.97	0.96	124
accuracy			0.89	500
macro avg	0.89	0.89	0.89	500
weighted avg	0.89	0.89	0.89	500